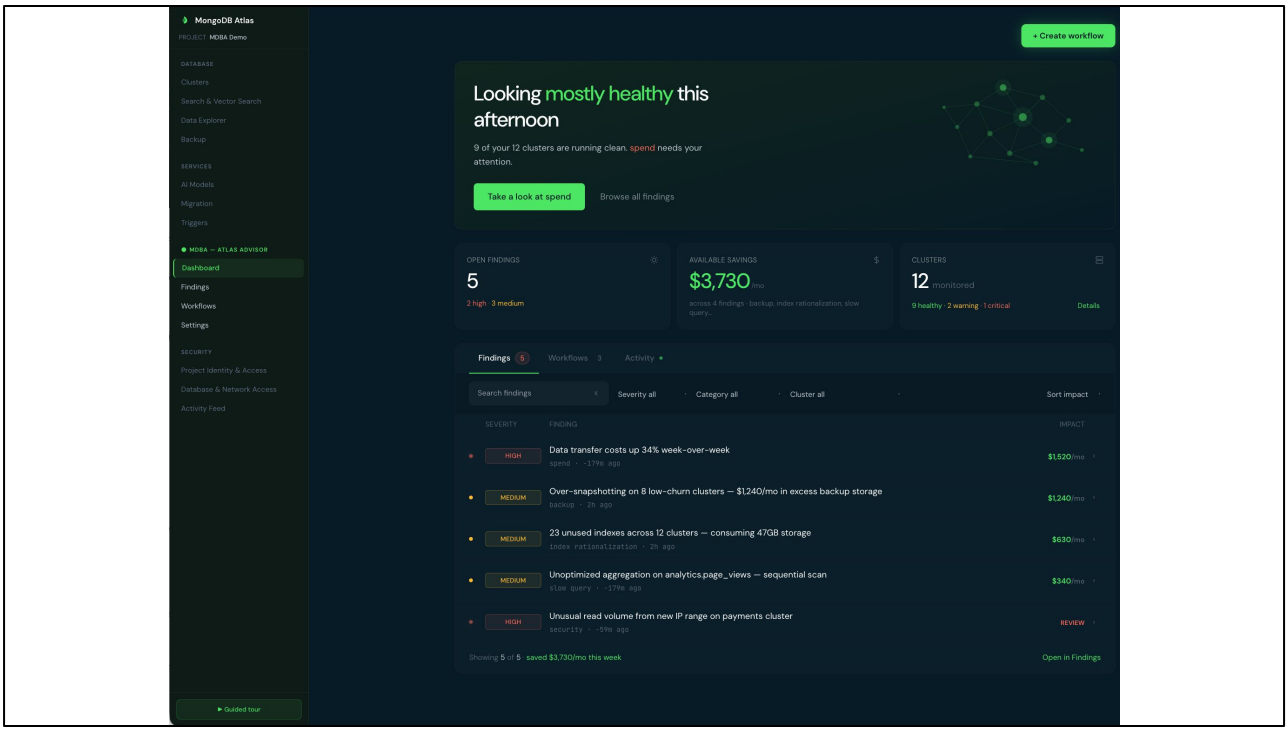

"This is the MDB Agent dashboard — this is what a customer sees first thing in the morning. And notice where it lives — right inside Atlas or on-prem — visible to the folks who have the power to act on it.

- ***Click on guided tour - get to At a Glance***

There are Three numbers tell the whole story:

- ***5 open findings** — 2 high, 3 medium*
- ***Estimated dollar savings per month** — in this case, $3,700 a month. That's $45K a year left on the table.*
- ***12 clusters monitored** — 9 healthy, 2 warning, 1 critical*

- ***Click next - Your workspace***

And if we look at the findings list — this is where one screen immediately provides full transparency on spend.

You can see exactly what type of savings each finding maps to: data transfer up 34% costing $1,520 a month, over-snapshotting on 8 clusters

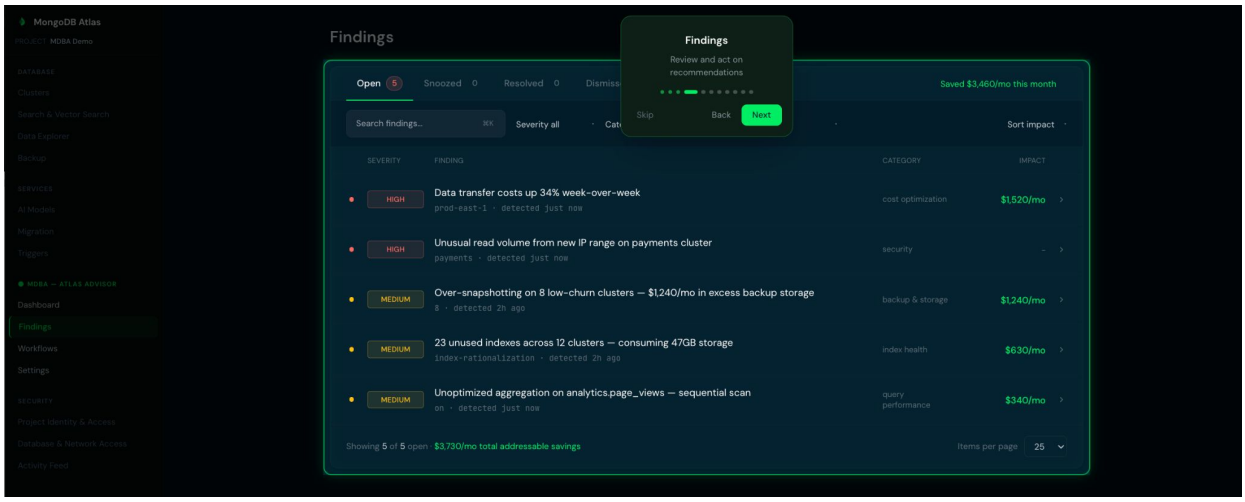

"If we wanted to drill down further — we can take a look at the high-impact items. Let's click into this top one, data transfer costs up 34%, and see what the agent actually found."

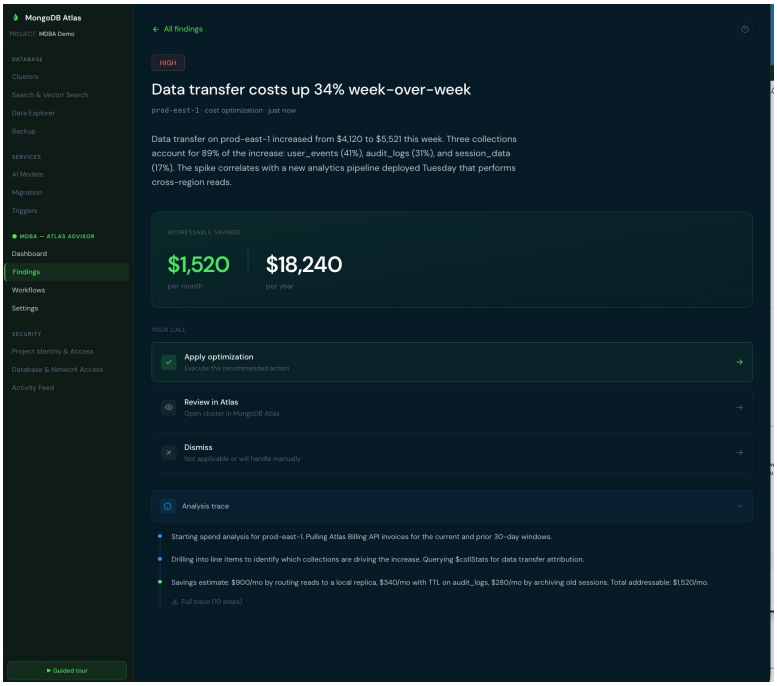

"So at this point you're getting the idea — something is driving these logic gates to provide proactive recommendations and make it obvious where the money is going.

Look at what the agent surfaced: data transfer on prod-east-1 jumped from $4,120 to $5,521 this week. It identified the three collections responsible — user_events, audit_logs, session_data — and traced the spike to a new analytics pipeline deployed Tuesday that performs cross-region reads.

Down here — addressable savings: $1,520 a month, $18,240 a year. And below that, the analysis trace — step by step, how the agent got there. It pulled the billing API, drilled into line items, queried $collStats, and arrived at three specific fixes with dollar estimates for each."

It also would support the full trace that you can download and view further.

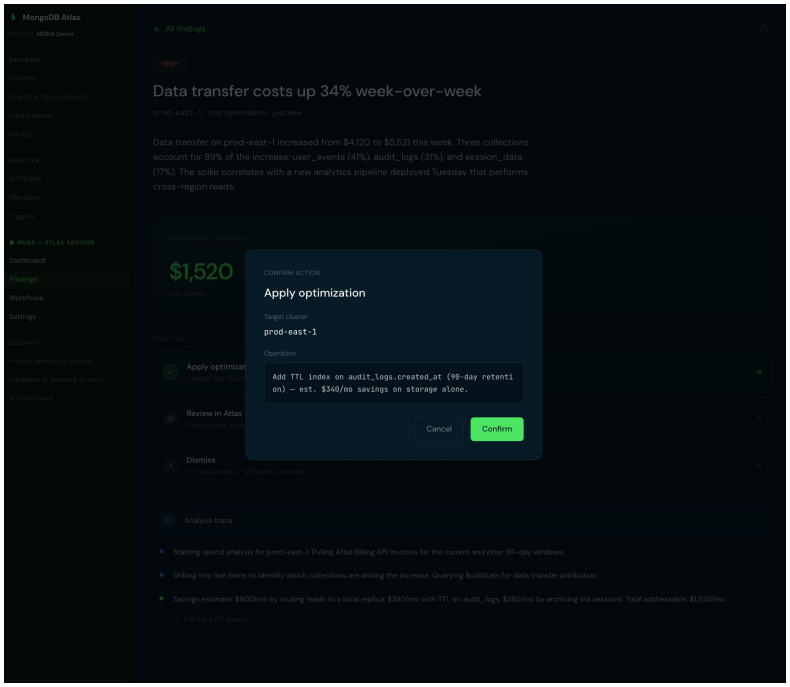

"And once we decide on the optimization we want to apply — this is the human-in-the-loop gate. It shows you the target cluster, the exact operation — in this case, add a TTL index on audit_logs.created_at with 90-day retention, estimated $340 a month in storage savings — and you confirm or cancel. Nothing touches the cluster without explicit approval."

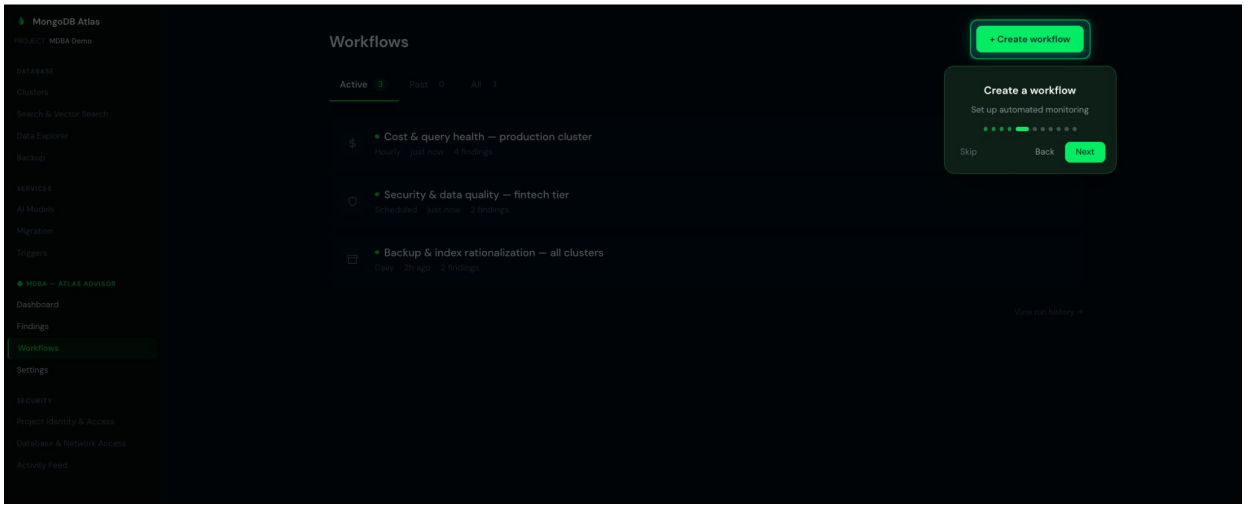

"Now — that finding came from somewhere. This is the Workflows pane. This is where customers can see what agents are actively deployed and watching their clusters. Think of this as triggers on steroids — cost & query health on the production cluster, security & data quality on the fintech tier, backup & index rationalization across all clusters.

These are running continuously. And if we want to build a new one — that's where it gets interesting."

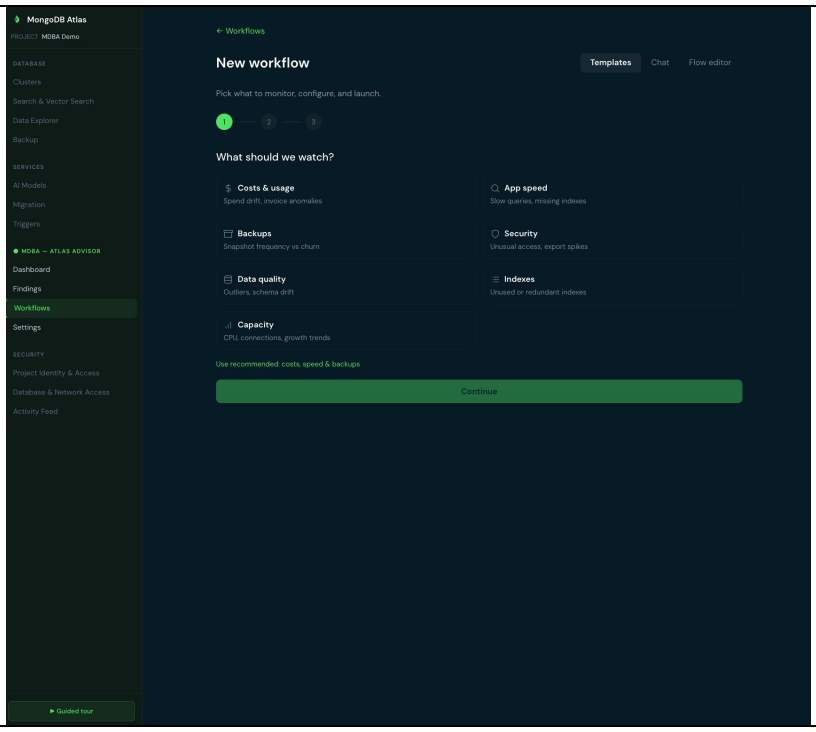

"Out of the box, you get the usual suspects — costs & usage, app speed, backups, security, data quality, indexes, capacity. The things every customer cares about, ready to deploy in one click.

But you also have two other paths. A chatbot interface — describe what you want to monitor in plain English and it builds the pipeline for you. Or a flow editor — a full visual sandbox where you wire up Atlas API calls, MongoDB operations, and LLM analysis nodes yourself. Three on-ramps for three personas."

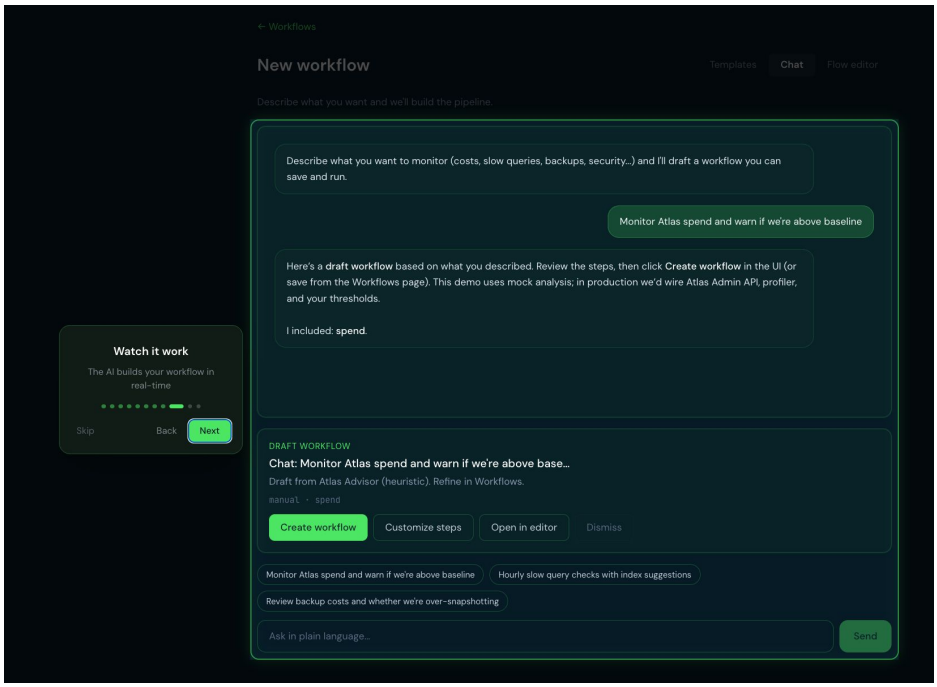

"Here's the chat path. The user typed 'Monitor Atlas spend and warn if we're above baseline' — plain English. The system drafted a workflow from that one sentence. It identified the category — spend — and generated a pipeline you can create immediately, customize the steps, or open in the full editor.

And notice the suggested prompts at the bottom — 'Hourly slow query checks with index suggestions,' 'Review backup costs and whether we're over-snapshotting.' These are the conversations customers are already having with their SAs. Now they can just type it."

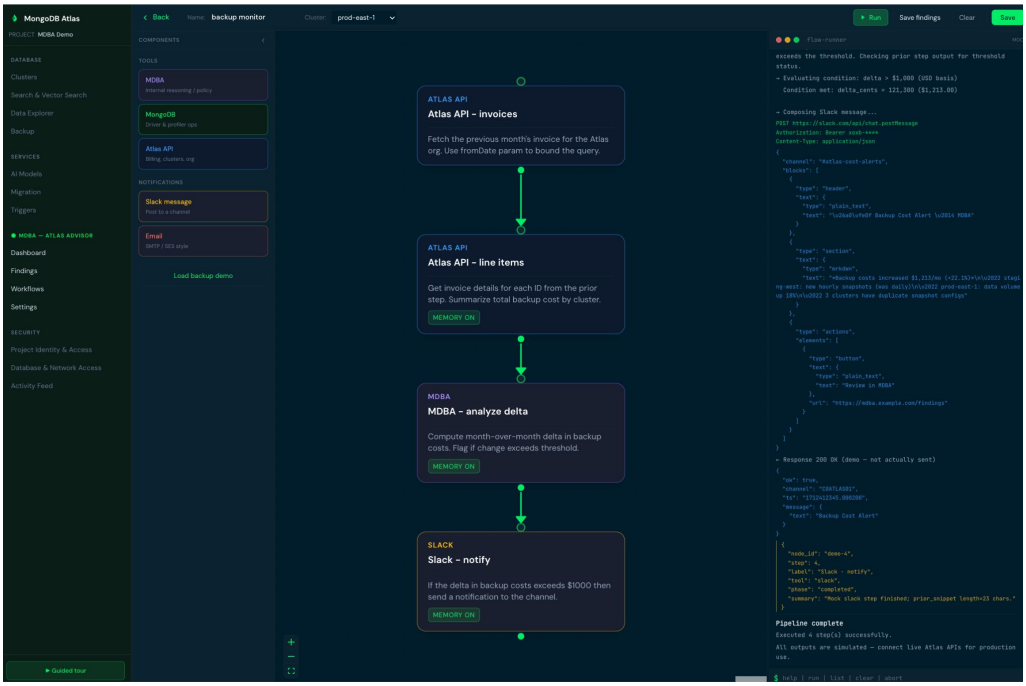

And for the folks who want full control — this is the flow editor. A visual DAG builder where you compose your own pipeline node by node.

What you're looking at is a backup monitor: Atlas API pulls the invoices, a second call gets the line items, MDBA analyzes the month-over-month delta, and if backup costs exceed the threshold — Slack notification. Four nodes, wired together, each one with memory on so the agent retains context between steps.

On the left — the component palette. MDBA, MongoDB, Atlas API, Slack, Email. Drag and drop. On the right — the terminal showing the actual pipeline execution, step by step. You can see the condition being evaluated, the Slack message being composed, the full trace. Complete observability."